

Double-click (or enter) to edit

Project: Machine Learning of Salary and Demographic Factors Name: Shachua Feng Supervisor:

Double-click (or enter) to edit

```
1 from google.colab import drive
2
3 # Mount Google Drive
4 drive.mount('/content/drive')

Mounted at /content/drive

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt

1 from sklearn.metrics import confusion_matrix
2 #from sklearn.metrics import plot_confusion_matrix

1 # from sklearn.metrics import plot_confusion_matrix doesn't work, so
2 !pip install plot_confusion_matrix

Collecting plot_confusion_matrix
  Downloading plot_confusion_matrix-0.0.2-py3-none-any.whl (3.6 kB)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from plot_confusion_matrix) (3.7.1)
Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (from plot_confusion_matrix) (1.23.5)
Requirement already satisfied: contourpy==1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->plot_confusion_matrix) (1.2.0)
Requirement already satisfied: cycler==0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->plot_confusion_matrix) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->plot_confusion_matrix) (4.44.3)
Requirement already satisfied: kiwisolver==1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->plot_confusion_matrix) (1.4.5)
Requirement already satisfied: packaging==20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->plot_confusion_matrix) (23.2)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->plot_confusion_matrix) (9.4.0)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->plot_confusion_matrix) (3.1.1)
Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->plot_confusion_matrix) (2.8.2)
Installing collected packages: plot_confusion_matrix
Successfully installed plot_confusion_matrix-0.0.2
```



```
1 #####
2 ## Read data and data wrangling
3 #####
4 # read in data loaded in google drive
5 file_path_1 = '/content/drive/My Drive/adult.data'
6 adult_1= pd.read_csv(file_path_1,header=None)
7 file_path_2 = '/content/drive/My Drive/adult.test.txt'
8 adult_2= pd.read_csv(file_path_2,header=None)
9 adult=pd.concat([adult_1, adult_2], ignore_index=True)

1 # add column names
2 cols=['age', 'workclass', 'fnlwt', 'education', 'education-num', 'marital-status', 'occupation', 'relationship', 'race', 'sex', 'capital-gain', 'capital-loss', 'hours-per-week', 'target']
3 adult.columns=cols
4
5 # add y column to data frame. target=1 for label '>50k' and y=0 for label '<=50k'
6 adult['target']=np.where(adult['label']=='>50k',1,0)
7 #adult['target']= adult['target'].astype(bool)
8 #
9 print(adult.describe())
10 adult.dtypes
11 adult.info()
```

	count	age	fnlwt	education-num	capital-gain	capital-loss	target
mean	38.643585	1.896641e+05	10.078089	1079.867626	87.502314	0.000000	0.160538
std	13.710510	1.056040e+05	2.570973	7452.019058	403.004552	0.000000	0.367108
min	17.000000	1.228500e+04	1.000000	0.000000	0.000000	0.000000	0.000000
25%	28.000000	1.175505e+05	9.000000	0.000000	0.000000	0.000000	0.000000
50%	37.000000	1.781445e+05	10.000000	0.000000	0.000000	0.000000	0.000000
75%	48.000000	2.376420e+05	12.000000	0.000000	0.000000	0.000000	0.000000
max	90.000000	1.490400e+06	16.000000	99999.000000	4356.000000	0.000000	1.000000

<Class 'pandas.core.frame.DataFrame'>
RangeIndex: 48842 entries, 0 to 48841
Data columns (total 16 columns):
Column Non-Null Count Dtype

age 48842 non-null int64

```
1 workclass 48842 non-null object
2 fnlwgt 48842 non-null int64
3 education 48842 non-null object
4 education-num 48842 non-null int64
5 marital-status 48842 non-null object
6 occupation 48842 non-null object
7 relationship 48842 non-null object
8 race 48842 non-null object
9 sex 48842 non-null object
10 capital-gain 48842 non-null int64
11 capital-loss 48842 non-null int64
12 hours-per-week 48842 non-null int64
13 native-country 48842 non-null object
14 label 48842 non-null object
15 target 48842 non-null int64
dtypes: int64(7), object(9)
memory usage: 6.0+ MB

1 # Data Manipulation: replace '?' with None
2 adult['workclass']=adult['workclass'].replace('?',None)
3 adult['occupation']=adult['occupation'].replace('?',None)
4 adult['native-country']=adult['native-country'].replace('?',None)
```

```
1 # Characterategorical columns
2 cols_cat=['workclass','education','marital-status','occupation','relationship','race','sex','native-country']
3
4 for x in cols_cat:
5     adult[x] = adult[x].astype('category')
6     #print(x)
7
8 adult.dtypes
```

```
age          int64
workclass    category
fnlwgt       int64
education    category
education-num int64
marital-status category
occupation   category
relationship category
race         category
sex          category
capital-gain int64
capital-loss int64
hours-per-week int64
native-country category
label        object
```

<https://colab.research.google.com/drive/1gMUOgOE4d8dWvSWXsuntrq49pb1qw#scrollTo=PqD5HSDkXD&printMode=true>

target
dtype: object

int64

1 # delete missing value

2 adult_cleaned=adult.dropna()

3 print(len(adult_cleaned))

45222

1 # drop original label

2 del adult_cleaned['label']

3 adult_cleaned.head(5)

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	NoI-in-family	White	Male	2174	0
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	NoI-in-family	White	Male	0	0
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0

<

>

```
1 #####
2 # Functions used by the machine learning algorithms
3 #Normalizing numeric data
4 def normalize(x):
5     if x.dtype == 'int' or x.dtype == 'float':
6         return ((x - min(x)) / (max(x) - min(x)))
7     else:
8         return x
```

<https://colab.research.google.com/drive/1gMUOgOE4d8dWvSWXsuntrq49pb1qw#scrollTo=PqD5HSDkXD&printMode=true>

```
11/20/23, 7:25 PM
1 #Converting categorical data to dummy/one-hot variables
2 def dummy(x):
3     cat_col='workclass','education','marital-status','occupation','relationship','race','sex','native-country']
4     x = pd.get_dummies(x, columns=cat_col, prefix = cat_col)
5     return x
6 #print the dataset
7 #adult_new.head(5)
8
9
10
11 # Print out Accuracy
12 def printAcc(y_test,y_pred):
13     from sklearn.metrics import accuracy_score
14
15     accuracy = accuracy_score(y_test, y_pred)
16     print(f"Accuracy: {accuracy}")
17
18
19 # print out confusion matrix
20 def printConfusion(y_test, y_pred):
21
22     from sklearn.metrics import confusion_matrix
23
24     cf=confusion_matrix(y_test, y_pred)
25     print(cf)
26     tn, fp, fn, tp=cf.ravel()
27     print ("Tp: ", tp, ", FP: ", fp, ", FN: ", fn)
```

<https://colab.research.google.com/drive/1gMUOgOE4d48duwYSWxsunlrq49pb1qw#scrollTo=PqD5HSDkXD8&printMode=true>

```
11/20/23, 7:25 PM
1 #####
2 # Decision tree
3 #####
4
5 # Decision tree without oversampling and normalization
6 # I want to use it as baseline to proof that oversampling and normalization improves decision tree
7 from sklearn.tree import DecisionTreeClassifier
8 from sklearn.model_selection import train_test_split
9
10 # Create a deep copy of the data frame adult_cleaned and name it adult_new
11 # adult_new is not normalized or oversampled
12 # I will use adult_new as the base line
13 adult_new=adult_cleaned.copy(deep=True)
14
15 #X = list(set(list(adult_cleaned)) - set(['target']))
16 x = adult_new.drop('target', axis=1)
17 y = adult_new['target']
18
19 # dummy variable
20 x=dummy(X)
21
22 # split train and test. test size is 0.35
23 X_train, X_test, y_train, y_test = train_test_split(x, y, test_size=0.35, random_state=42)
24
25 # create decision tree classifier
26 dtc = DecisionTreeClassifier(random_state=52)
27 # Train the model on the training data
28 dtc.fit(X_train, y_train)
29 # Make predictions on the test data
30 y_pred = dtc.predict(X_test)
31
32 print("Decision Tree for data without normalization and oversampling")
33 # Print Accuracy
34 printAcc(y_test,y_pred)
35
36 # Print Confusion Matrix
37 print("")
38 print("Confusion Matrix")
39 printConfusion(y_test, y_pred)
40
41 # Print Diagnosis
42 print("")
43 printReport(y_test,y_pred)
44
45
```

<https://colab.research.google.com/drive/1gMUOgOE4d48duwYSWxsunlrq49pb1qw#scrollTo=PqD5HSDkXD8&printMode=true>

Decision Tree for data without normalization and oversampling
Accuracy: 0.7998483699772555

Confusion Matrix

[[11507 1697]
 [1471 1153]]

TP: 1153 , FP: 1697 , TN: 11507 , FN: 1471

	precision	recall	f1-score	support
0	0.89	0.87	0.88	13204
1	0.40	0.44	0.42	2624
accuracy			0.80	15828
macro avg	0.65	0.66	0.65	15828
weighted avg	0.81	0.80	0.80	15828

11/20/23, 7:25 PM

820_adult_ML.ipynb - Colaboratory

```
1 # Decision tree with normalization but no oversampling
2
3 from sklearn.tree import DecisionTreeClassifier
4 from sklearn.model_selection import train_test_split
5
6 # Create a deep copy of the data frame adult_cleaned and name it adult_new
7 # adult_new is not normalized or oversampled
8 # I will use adult_new as the base line
9 adult_dt_norm = adult_cleaned.copy(deep=True)
10 x = adult_dt_norm.drop('target', axis=1)
11 y = adult_dt_norm['target']
12
13 # Normalization the numerical columns
14 num_cols = ['age', 'fnlwt', 'education-num', 'capital-gain', 'capital-loss', 'hours-per-week']
15 # why this code doesn't work?
16 #X[num_cols] = X[num_cols].apply(lambda x:normalize(x), axis=0)
17 X_num = X[num_cols]
18 X_num_normalized = X_num.apply(normalize, axis=0)
19 # combine the normalized numerical columns with the categorical columns
20 X = pd.concat([X_num_normalized, X.drop(num_cols, axis=1)]), axis=1)
21
22 #print(X.head(5))
23 #print(X.tail(5))
24 # dummy variable
25 X=dummy(X)
26 print(X.head(5))
27 # split train and test. test size is 0.35
28 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.35, random_state=52)
29
30 # create decision tree classifier
31 dtc_norm = DecisionTreeClassifier(random_state=52)
32 # Train the model on the training data
33 dtc_norm.fit(X_train, y_train)
34 # Make predictions on the test data
35 y_pred = dtc_norm.predict(X_test)
36
37 print("Decision Tree for data with normalization but no oversampling")
38 # Print Accuracy
39 printAcc(y_test,y_pred)
40
41 # Print Confusion Matrix
42 print("")
43 print("Confusion Matrix")
44 printConfusion(y_test, y_pred)
45
46 # Print Diagnosis
```

```
47 print("")
48 printReport(y_test,y_pred)
49
50
51
52
53
```

```

age      fnlwgt  education-num  capital-gain  capital-loss  \
0  0.301370  0.043350  0.800000  0.02174  0.0
1  0.452055  0.047274  0.800000  0.00000  0.0
2  0.287671  0.136877  0.533333  0.00000  0.0
3  0.493151  0.149792  0.400000  0.00000  0.0
4  0.150685  0.219998  0.800000  0.00000  0.0

hours-per-week  workclass_ Federal-gov  workclass_ Local-gov  \
0  0.397959  0  0
1  0.122449  0  0
2  0.397959  0  0
3  0.397959  0  0
4  0.397959  0  0

workclass_ Never-worked  workclass_ Private  ...  native-country_ Portugal  \
0  0  0  0 ...  0
1  0  0 ...  0
2  0  1 ...  0
3  0  1 ...  0
4  0  1 ...  0

native-country_ Puerto-Rico  native-country_ Scotland  \
0  0  0
1  0  0
2  0  0
3  0  0
4  0  0

native-country_ South  native-country_ Taiwan  native-country_ Thailand  \
0  0  0  0
1  0  0  0
2  0  0  0
3  0  0  0
4  0  0  0

native-country_ Trinidad&Tobago  native-country_ United-States  \
0  0  1
1  0  1
2  0  1
3  0  1
4  0  0
```

<https://colab.research.google.com/drive/1gMUOgOE4d8dWwSWXsuntrq49pb1qw#scrollTo=PqD5HSDkXDe&printMode=true>

```

native-country_ Vietnam  native-country_ Yugoslavia
0  0  0
1  0  0
2  0  0
3  0  0
4  0  0

[5 rows x 105 columns]
Decision Tree for data with normalization but no oversampling
Accuracy: 0.8035759413697245

Confusion Matrix
[[11590 1636]
 [ 1473 1129]]
TP: 1129 , FP: 1636 , TN: 11590 , FN: 1473
```

```
1 # Decision tree with normalization and random oversampling
2 from sklearn.tree import DecisionTreeClassifier
3 from sklearn.model_selection import train_test_split
4 from imblearn.over_sampling import RandomOverSampler
5
6 # Create a deep copy of the data frame adult_dt_norm_ros
7 # X is normalized and random oversampled
8 adult_dt_norm_ros = adult_cleaned.copy(deep=True)
9 X = adult_dt_norm_ros.drop('target', axis=1)
10 y = adult_dt_norm_ros['target']
11
12 # Normalization the numerical columns
13 num_cols = ['age', 'fhlwgt', 'education-num', 'capital-gain', 'capital-loss', 'hours-per-week']
14 X_num = X[num_cols]
15 X_num.normalized = X_num.apply(normalize, axis=0)
16 # combine the normalized numerical columns with the categorical columns
17 X = pd.concat([X_num.normalized, X.drop(num_cols, axis=1)], axis=1)
18
19 # dummy variable
20 X=dummy(X)
21
22 # split train and test. test size is 0.35
23 X_train_norm, X_test_norm, y_train_norm, y_test_norm = train_test_split(X, y, test_size=0.35, random_state=52)
24
25 #Random Oversampling
26 ros = RandomOverSampler(sampling_strategy='auto', random_state=52)
27 X_train_ros, y_train_ros = ros.fit_resample(X_train_norm, y_train_norm)
28
29 # create decision tree classifier for normalized data with random oversampling
30 dtc_norm_ros = DecisionTreeClassifier(random_state=52)
31 # Train the model on the training data
32 dtc_norm.fit(X_train_ros, y_train_ros)
33 # Make predictions on the test data
34 y_pred_norm_ros = dtc_norm.predict(X_test_norm)
35
36 print("Decision Tree for data with normalization and random oversampling")
37 # Print Accuracy
38 printAcc(y_test_norm,y_pred_norm_ros)
39
40 # Print Confusion Matrix
41 print("")
42 print("Confusion Matrix")
43 printConfusion(y_test_norm, y_pred_norm_ros)
44
45 # Print Diagnosis
```

https://colab.research.google.com/drive/1gMUOgOE4d8duwSWXsuntrq49pb1qw#scrollTo=PqD5HSDkXD&printMode=true

```
46 print("")
47 printReport(y_test_norm,y_pred_norm_ros)

Decision Tree for data with normalization and random oversampling
Accuracy: 0.803702297220116

Confusion Matrix
[[11656 1570]
 [ 1537 1065]]
TP: 1065 , FP: 1570 , TN: 11656 , FN: 1537

precision    recall  f1-score   support

   0    0.88    0.88    0.88    13226
   1    0.40    0.41    0.41     2602

 accuracy
macro avg    0.64    0.65    0.64    15828
weighted avg    0.80    0.80    0.80    15828
```

Double-click (or enter) to edit

11/20/23, 7:25 PM820_adult_ML.ipynb - Colaboratory

```
1 # Cross validated decision tree normalized no over sampling
2 from sklearn.model_selection import cross_val_score, cross_val_predict, KFold
3 # StratifiedKFold
4 from sklearn.tree import DecisionTreeClassifier
5 from imblearn.pipeline import Pipeline
6 from sklearn.metrics import classification_report
7
8 # Create a deep copy of the data frame adult_cleaned and name it adult_c_dt_norm_ros
9 X is normalized and random oversampled
10 adult_c_dt = adult_cleaned.copy(deep=True)
11 X = adult_c_dt.drop('target', axis=1)
12 y = adult_c_dt['target']
13
14 # Normalization the numerical columns
15 num_cols = ['age', 'fmlwgt', 'education-num', 'capital-gain', 'capital-loss', 'hours-per-week']
16 X_num = X[num_cols]
17 X_num.normalized = X_num.apply(normalize, axis=0)
18 # combine the normalized numerical columns with the categorical columns
19 X = pd.concat([X_num.normalized, X.drop(num_cols, axis=1)], axis=1)
20
21 # dummy variable
22 X=dummy(X)
23
24 # Create a Decision Tree classifier
25 c_dt= DecisionTreeClassifier()
26
27 # Create a pipeline with oversampling and the decision tree classifier
28 model = Pipeline([('ros', ros), ('dt', c_dt_rom)])
29
30 # Set up cross-validation using StratifiedKFold
31 cv = KFold(n_splits=10, shuffle=True, random_state=52)
32
33 # Perform cross-validation and obtain predicted labels
34 y_pred = cross_val_predict(c_dt,X,y, cv=cv)
35
36 # Calculate and print classification report
37 print("Classification Report:\n", classification_report(y, y_pred))
38
39 # Perform cross-validation
40 scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv)
41
42 # Print the mean accuracy across all folds
43 print("Mean Accuracy:", scores.mean())
44
```

<https://colab.research.google.com/drive/1gMUOgOE4d8duwYSWxsunlrq49pb1qw#scrollTo=PqD5HSDkXD8&printMode=true>

11/20/23, 7:25 PM820_adult_ML.ipynb - Colaboratory

```
1 # cross validated decision tree normalized random oversampling
2 # I use pipeline!
3 from sklearn.model_selection import cross_val_score, cross_val_predict, KFold
4 # StratifiedKFold
5 from sklearn.tree import DecisionTreeClassifier
6 from imblearn.over_sampling import RandomOverSampler
7 from imblearn.pipeline import Pipeline
8 from sklearn.metrics import classification_report
9 import pandas as pd
10
11 # Create a deep copy of the data frame adult_cleaned and name it adult_c_dt_norm_ros
12 X is normalized and random oversampled
13 adult_c_dt_norm_ros = adult_cleaned.copy(deep=True)
14 X = adult_c_dt_norm_ros.drop('target', axis=1)
15 y = adult_c_dt_norm_ros['target']
16
17 # Normalization the numerical columns
18 num_cols = ['age', 'fmlwgt', 'education-num', 'capital-gain', 'capital-loss', 'hours-per-week']
19 X_num = X[num_cols]
20 X_num.normalized = X_num.apply(normalize, axis=0)
21 # combine the normalized numerical columns with the categorical columns
22 X = pd.concat([X_num.normalized, X.drop(num_cols, axis=1)], axis=1)
23
24 # dummy variable
25 X=dummy(X)
26
27 # Create a RandomOverSampler
28 ros = RandomOverSampler()
29
30 # Create a Decision Tree classifier
31 c_dt_rom = DecisionTreeClassifier()
32
33 # Apply oversampling to X and y
34 X_resampled, y_resampled = ros.fit_resample(X, y)
35
```

<https://colab.research.google.com/drive/1gMUOgOE4d8duwYSWxsunlrq49pb1qw#scrollTo=PqD5HSDkXD8&printMode=true>

```
35
36
37 # Create a pipeline with oversampling and the decision tree classifier
38 model = Pipeline([('ros', ros), ('dt', c_dt_om)])
39
40 # Set up cross-validation using StratifiedKFold
41 cv = KFold(n_splits=10, shuffle=True, random_state=52)
42
43 # Perform cross-validation and obtain predicted labels
44 y_pred = cross_val_predict(model, X_resampled, y_resampled, cv=cv)
45
46 # Calculate and print classification report
47 print("Classification Report:\n", classification_report(y_resampled, y_pred))
48
49 # Perform cross-validation
50 scores = cross_val_score(model, X, y, scoring='accuracy', cv=cv)
51
52 # Print the mean accuracy across all folds
53 print("Mean Accuracy:", scores.mean())
54
55 ##### comments: how to print out all those stuff?
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.88	0.93	37714
1	0.89	1.00	0.94	37714
accuracy			0.94	75428
macro avg	0.94	0.94	0.94	75428
weighted avg	0.94	0.94	0.94	75428

Mean Accuracy: 0.8071071215644292

```
1 #####
2 # Logistic regression
3 #####
4
5 # Logistic regression with cross validation and no oversampling
6 from sklearn.linear_model import LogisticRegression
7 from sklearn.model_selection import cross_val_score, KFold
8
9 # Create a deep copy of the data frame adult_cleaned and name it adult_c_dt_norm_ros
10 # X is normalized and random oversampled
11 adult_c_log = adult_cleaned.copy(deep=True)
12 X = adult_c_log.drop('target', axis=1)
13 y = adult_c_log['target']
14
15 # Normalization the numerical columns
16 num_cols = ['age', 'fmlwgt', 'education-num', 'capital-gain', 'capital-loss', 'hours-per-week']
17 X_num = X[num_cols]
18 X_num_normalized = X_num.apply(normalize, axis=0)
19 # combine the normalized numerical columns with the categorical columns
20 X = pd.concat([X_num_normalized, X.drop(num_cols, axis=1)], axis=1)
21
22 # dummy variable
23 X=dummy(X)
24
25 # create a logistic model
26 model = LogisticRegression(max_iter=1000) # Increase max_iter if needed for convergence
27
28 # Set up 10-fold cross-validation
29 kfold = KFold(n_splits=10, shuffle=True, random_state=52)
30
31 # Perform cross-validation and get the accuracy scores for each fold
32 scores = cross_val_score(model, X, y, cv=kfold)
33
34 # Print the accuracy for each fold and the mean accuracy
35 for i, score in enumerate(scores, 1):
36     print(f'Fold {i}: {score}')
```

```
37
38 print(f'Mean Accuracy: {scores.mean()}')
```

Fold 1: 0.8425823568428035
Fold 2: 0.853415874419633
Fold 3: 0.8423264042459089
Fold 4: 0.8491817779743477
Fold 5: 0.8524988942945599
Fold 6: 0.842998275099513
Fold 7: 0.85161443299425033

Fold 8: 0.8398938522777533
Fold 9: 0.8403361344537815
Fold 10: 0.8423264042459089
Mean Accuracy: 0.8457165856207152

<https://colab.research.google.com/drive/1gMUOgOE4d8duwYSWxsunlrq49pb1qw#scrollTo=PqD5HSDkXD8&printMode=true>

```
1 # Logistic regression with cross validation and random oversampling
2
3 from sklearn.linear_model import LogisticRegression
4 from sklearn.model_selection import cross_val_score, KFold
5 from imblearn.over_sampling import RandomOverSampler
6
7 # Create a deep copy of the data frame adult_cleaned and name it adult_c_dt_norm_ros
8 # X is normalized and random oversampled
9 adult_c_log_ros = adult_cleaned.copy(deep=True)
10 X = adult_c_log_ros.drop('target', axis=1)
11 y = adult_c_log_ros['target']
12
13 # Normalization the numerical columns
14 num_cols = ['age', 'fnlwt', 'education-num', 'capital-gain', 'capital-loss', 'hours-per-week']
15 X_num = X[num_cols]
16 X_num_normalized = X_num.apply(normalize, axis=0)
17 # combine the normalized numerical columns with the categorical columns
18 X = pd.concat([X_num_normalized, X.drop(num_cols, axis=1)], axis=1)
19
20 # dummy
21 X=dummy(X)
22
```

<https://colab.research.google.com/drive/1gMUOgOE4d8duwYSWxsunlrq49pb1qw#scrollTo=PqD5HSDkXD8&printMode=true>